

pydantic_ai.models

Logic related to making requests to an LLM.

The aim here is to make a common interface for different LLMs, so that the rest of the code can be agnostic to the specific LLM being used.

KnownModelName module-attribute

```
KnownModelName = Literal[
    "openai:gpt-4o",
    "openai:gpt-4o-mini",
    "openai:gpt-4-turbo",
    "openai:gpt-4",
    "openai:o1-preview",
    "openai:o1-mini",
    "openai:gpt-3.5-turbo",
    "groq:llama3-1-70b-versatile",
    "groq:llama3-groq-70b-8192-tool-use-preview",
    "groq:llama3-groq-8b-8192-tool-use-preview",
    "groq:llama3-1-70b-specdec",
    "groq:llama3-1-8b-instant",
    "groq:llama3-2-1b-preview",
    "groq:llama3-2-3b-preview",
    "groq:llama3-2-11b-vision-preview",
    "groq:llama3-2-90b-vision-preview",
    "groq:llama3-70b-8192",
    "groq:llama3-8b-8192",
    "groq:mixtral-8x7b-32768",
    "groq:gemma2-9b-it",
    "groq:gemma-7b-it",
    "gemini-1.5-flash",
    "gemini-1.5-pro",
    "vertexai:gemini-1.5-flash",
    "vertexai:gemini-1.5-pro",
    "ollama:codellama",
    "ollama:gemma",
    "ollama:gemma2",
    "ollama:llama3",
    "ollama:llama3.1",
    "ollama:llama3.2",
    "ollama:llama3.2-vision",
    "ollama:llama3.3",
    "ollama:mistral",
    "ollama:mistral-nemo",
    "ollama:mixtral",
    "ollama:phi3",
    "ollama:qwen",
    "ollama:qwen2",
    "ollama:qwen2.5",
    "ollama:starcode2",
    "test",
]
```

Known model names that can be used with the `model` parameter of `Agent`.

`KnownModelName` is provided as a concise way to specify a model.

Model

Bases: `ABC`

Abstract class for a model.

```
Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py

77 class Model(ABC):
78     """Abstract class for a model."""
79
80     @abstractmethod
81     async def agent_model(
82         self,
83         *,
84         function_tools: list[ToolDefinition],
85         allow_text_result: bool,
86         result_tools: list[ToolDefinition],
87     ) -> AgentModel:
88         """Create an agent model, this is called for each step of an agent run.
89
90         This is async in case slow/async config checks need to be performed that can't be done in '__init__'.
91
92         Args:
93             function_tools: The tools available to the agent.
94             allow_text_result: Whether a plain text final response/result is permitted.
95             result_tools: Tool definitions for the final result tool(s), if any.
96
97         Returns:
98             An agent model.
99         """
100         raise NotImplementedError()
101
102     @abstractmethod
103     def name(self) -> str:
104         raise NotImplementedError()
```

agent_model abstractmethod async

```
agent_model(
    *,
    function_tools: list[ToolDefinition],
    allow_text_result: bool,
    result_tools: list[ToolDefinition]
) -> AgentModel
```

Create an agent model, this is called for each step of an agent run.

This is async in case slow/async config checks need to be performed that can't be done in `__init__`.

Parameters:

Name	Type	Description	Default
<code>function_tools</code>	<code>list[ToolDefinition]</code>	The tools available to the agent.	<i>required</i>
<code>allow_text_result</code>	<code>bool</code>	Whether a plain text final response/result is permitted.	<i>required</i>
<code>result_tools</code>	<code>list[ToolDefinition]</code>	Tool definitions for the final result tool(s), if any.	<i>required</i>

Returns:

Type	Description
<code>AgentModel</code>	An agent model.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/__init__.py
80 @abstractmethod
81 async def agent_model(
82     self,
83     *,
84     function_tools: list[ToolDefinition],
85     allow_text_result: bool,
86     result_tools: list[ToolDefinition],
87 ) -> AgentModel:
88     """Create an agent model, this is called for each step of an agent run.
89
90     This is async in case slow/async config checks need to be performed that can't be done in '__init__'.
91
92     Args:
93         function_tools: The tools available to the agent.
94         allow_text_result: Whether a plain text final response/result is permitted.
95         result_tools: Tool definitions for the final result tool(s), if any.
96
97     Returns:
98         An agent model.
99     """
100     raise NotImplementedError()
```

AgentModel

Bases: `ABC`

Model configured for each step of an Agent run.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/__init__.py
107 class AgentModel(ABC):
108     """Model configured for each step of an Agent run."""
109
110     @abstractmethod
111     async def request(self, messages: list[Message]) -> tuple[ModelAnyResponse, Cost]:
112         """Make a request to the model."""
113         raise NotImplementedError()
114
115     @asynccontextmanager
116     async def request_stream(self, messages: list[Message]) -> AsyncIterator[EitherStreamedResponse]:
117         """Make a request to the model and return a streaming response."""
118         raise NotImplementedError(f'Streamed requests not supported by this {self.__class__.__name__}')
119         # yield is required to make this a generator for type checking
120         # noinspection PyUnreachableCode
121         yield # pragma: no cover
```

`request` `abstractmethod` `async`

```
request(
    messages: list[Message],
) -> tuple[ModelAnyResponse, Cost]
```

Make a request to the model.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/__init__.py
110 @abstractmethod
111 async def request(self, messages: list[Message]) -> tuple[ModelAnyResponse, Cost]:
112     """Make a request to the model."""
113     raise NotImplementedError()
```

`request_stream` `async`

```
request_stream(
    messages: list[Message],
) -> AsyncIterator[EitherStreamedResponse]
```

Make a request to the model and return a streaming response.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/__init__.py
115 @asynccontextmanager
116 async def request_stream(self, messages: list[Message]) -> AsyncIterator[EitherStreamedResponse]:
117     """Make a request to the model and return a streaming response."""
118     raise NotImplementedError(f'Streamed requests not supported by this {self.__class__.__name__}')
119     # yield is required to make this a generator for type checking
120     # noinspection PyUnreachableCode
121     yield # pragma: no cover
```

StreamTextResponse

Streamed response from an LLM when returning text.

Source code in `pydantic_ai_slim/pydantic_ai/models/__init__.py`

```
124 class StreamTextResponse(ABC):
125     """Streamed response from an LLM when returning text."""
126
127     def __aiter__(self) -> AsyncIterator[None]:
128         """Stream the response as an async iterable, building up the text as it goes.
129
130         This is an async iterator that yields 'None' to avoid doing the work of validating the input and
131         extracting the text field when it will often be thrown away.
132         """
133         return self
134
135     @abstractmethod
136     async def __anext__(self) -> None:
137         """Process the next chunk of the response, see above for why this returns 'None'."""
138         raise NotImplementedError()
139
140     @abstractmethod
141     def get(self, *, final: bool = False) -> Iterable[str]:
142         """Returns an iterable of text since the last call to 'get()' - e.g. the text delta.
143
144         Args:
145             final: If True, this is the final call, after iteration is complete, the response should be fully validated
146                     and all text extracted.
147         """
148         raise NotImplementedError()
149
150     @abstractmethod
151     def cost(self) -> Cost:
152         """Return the cost of the request.
153
154         NOTE: this won't return the full cost until the stream is finished.
155         """
156         raise NotImplementedError()
157
158     @abstractmethod
159     def timestamp(self) -> datetime:
160         """Get the timestamp of the response."""
161         raise NotImplementedError()
```

`__aiter__`

`__aiter__()` -> `AsyncIterator[None]`

Stream the response as an async iterable, building up the text as it goes.

This is an async iterator that yields `None` to avoid doing the work of validating the input and extracting the text field when it will often be thrown away.

Source code in `pydantic_ai_slim/pydantic_ai/models/__init__.py`

```
127 def __aiter__(self) -> AsyncIterator[None]:
128     """Stream the response as an async iterable, building up the text as it goes.
129
130     This is an async iterator that yields 'None' to avoid doing the work of validating the input and
131     extracting the text field when it will often be thrown away.
132     """
133     return self
```

`__anext__` abstractmethod async

`__anext__()` -> `None`

Process the next chunk of the response, see above for why this returns `None`.

Source code in `pydantic_ai_slim/pydantic_ai/models/__init__.py`

```
135 @abstractmethod
136 async def __anext__(self) -> None:
137     """Process the next chunk of the response, see above for why this returns 'None'."""
138     raise NotImplementedError()
```

`get` abstractmethod

`get(*, final: bool = False) -> Iterable[str]`

Returns an iterable of text since the last call to `get()` - e.g. the text delta.

Parameters:

Name	Type	Description	Default
<code>final</code>	<code>bool</code>	If True, this is the final call, after iteration is complete, the response should be fully validated and all text extracted.	<code>False</code>

Source code in `pydantic_ai_slim/pydantic_ai/models/__init__.py`

```
140 @abstractmethod
141 def get(self, *, final: bool = False) -> Iterable[str]:
142     """Returns an iterable of text since the last call to 'get()' - e.g. the text delta.
143
144     Args:
145         final: If True, this is the final call, after iteration is complete, the response should be fully validated
146                 and all text extracted.
147     """
148     raise NotImplementedError()
```

`cost` abstractmethod

```
cost() -> Cost
.....
```

Return the cost of the request.

NOTE: this won't return the full cost until the stream is finished.

99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py

```
150 @abstractmethod
151 def cost(self) -> Cost:
152     """Return the cost of the request.
153
154     NOTE: this won't return the full cost until the stream is finished.
155     """
156     raise NotImplementedError()
```

timestamp abstractmethod

```
timestamp() -> datetime
.....
```

Get the timestamp of the response.

99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py

```
158 @abstractmethod
159 def timestamp(self) -> datetime:
160     """Get the timestamp of the response."""
161     raise NotImplementedError()
```

StreamStructuredResponse

Bases: [ABC](#)

Streamed response from an LLM when calling a tool.

99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py

```
164 class StreamStructuredResponse(ABC):
165     """Streamed response from an LLM when calling a tool."""
166
167     def __aiter__(self) -> AsyncIterator[None]:
168         """Stream the response as an async iterable, building up the tool call as it goes.
169
170         This is an async iterator that yields 'None' to avoid doing the work of building the final tool call when
171         it will often be thrown away.
172         """
173         return self
174
175     @abstractmethod
176     async def __anext__(self) -> None:
177         """Process the next chunk of the response, see above for why this returns 'None'."""
178         raise NotImplementedError()
179
180     @abstractmethod
181     def get(self, *, final: bool = False) -> ModelStructuredResponse:
182         """Get the 'ModelStructuredResponse' at this point.
183
184         The 'ModelStructuredResponse' may or may not be complete, depending on whether the stream is finished.
185
186         Args:
187             final: If True, this is the final call, after iteration is complete, the response should be fully validated.
188         """
189         raise NotImplementedError()
190
191     @abstractmethod
192     def cost(self) -> Cost:
193         """Get the cost of the request.
194
195         NOTE: this won't return the full cost until the stream is finished.
196         """
197         raise NotImplementedError()
198
199     @abstractmethod
200     def timestamp(self) -> datetime:
201         """Get the timestamp of the response."""
202         raise NotImplementedError()
```

__aiter__

```
__aiter__() -> AsyncIterator[None]
.....
```

Stream the response as an async iterable, building up the tool call as it goes.

This is an async iterator that yields `None` to avoid doing the work of building the final tool call when it will often be thrown away.

99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py

```
167 def __aiter__(self) -> AsyncIterator[None]:
168     """Stream the response as an async iterable, building up the tool call as it goes.
169
170     This is an async iterator that yields 'None' to avoid doing the work of building the final tool call when
171     it will often be thrown away.
172     """
173     return self
```

__anext__ abstractmethod async

```
__anext__() -> None
.....
```

Process the next chunk of the response, see above for why this returns `None`.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py
175 @abstractmethod
176 async def __anext__(self) -> None:
177     """Process the next chunk of the response, see above for why this returns 'None'."""
178     raise NotImplementedError()
```

get abstractmethod

```
get(*, final: bool = False) -> ModelStructuredResponse
```

Get the `ModelStructuredResponse` at this point.

The `ModelStructuredResponse` may or may not be complete, depending on whether the stream is finished.

Parameters:

Name	Type	Description	Default
final	bool	If True, this is the final call, after iteration is complete, the response should be fully validated.	False

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py
180 @abstractmethod
181 def get(self, *, final: bool = False) -> ModelStructuredResponse:
182     """Get the 'ModelStructuredResponse' at this point.
183
184     The 'ModelStructuredResponse' may or may not be complete, depending on whether the stream is finished.
185
186     Args:
187         final: If True, this is the final call, after iteration is complete, the response should be fully validated.
188     """
189     raise NotImplementedError()
```

cost abstractmethod

```
cost() -> Cost
```

Get the cost of the request.

NOTE: this won't return the full cost until the stream is finished.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py
191 @abstractmethod
192 def cost(self) -> Cost:
193     """Get the cost of the request.
194
195     NOTE: this won't return the full cost until the stream is finished.
196     """
197     raise NotImplementedError()
```

timestamp abstractmethod

```
timestamp() -> datetime
```

Get the timestamp of the response.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/_init_.py
199 @abstractmethod
200 def timestamp(self) -> datetime:
201     """Get the timestamp of the response."""
202     raise NotImplementedError()
```

ALLOW_MODEL_REQUESTS module-attribute

```
ALLOW_MODEL_REQUESTS = True
```

Whether to allow requests to models.

This global setting allows you to disable request to most models, e.g. to make sure you don't accidentally make costly requests to a model during tests.

The testing models `TestModel` and `FunctionModel` are no affected by this setting.

check_allow_model_requests

```
check_allow_model_requests() -> None
```

Check if model requests are allowed.

If you're defining your own models that have cost or latency associated with their use, you should call this in `Model.agent_model`.

Raises:

Type	Description
<code>RuntimeError</code>	If model requests are not allowed.

```
219 def check_allow_model_requests() -> None:
220     """Check if model requests are allowed.
221
222     If you're defining your own models that have cost or latency associated with their use, you should call this in
223     ['Model.agent_model'] [pydantic_ai.models.Model.agent_model].
224
225     Raises:
226         RuntimeError: If model requests are not allowed.
227     """
228     if not ALLOW_MODEL_REQUESTS:
229         raise RuntimeError('Model requests are not allowed, since ALLOW_MODEL_REQUESTS is False')
```

override_allow_model_requests

```
override_allow_model_requests(
    allow_model_requests: bool,
) -> Iterator[None]
```

Context manager to temporarily override `ALLOW_MODEL_REQUESTS`.

Parameters:

Name	Type	Description	Default
allow_model_requests	bool	Whether to allow model requests within the context.	required

```
232 @contextmanager
233 def override_allow_model_requests(allow_model_requests: bool) -> Iterator[None]:
234     """Context manager to temporarily override ['ALLOW_MODEL_REQUESTS'] [pydantic_ai.models.ALLOW_MODEL_REQUESTS].
235
236     Args:
237         allow_model_requests: Whether to allow model requests within the context.
238     """
239     global ALLOW_MODEL_REQUESTS
240     old_value = ALLOW_MODEL_REQUESTS
241     ALLOW_MODEL_REQUESTS = allow_model_requests # pyright: ignore[reportConstantRedefinition]
242     try:
243         yield
244     finally:
245         ALLOW_MODEL_REQUESTS = old_value # pyright: ignore[reportConstantRedefinition]
```