

pydantic_ai.models.gemini

Custom interface to the `generativelanguage.googleapis.com` API using [HTTPX](#) and [\[Pydantic\]\(https://docs.pydantic.dev/latest/\)](#).

The Google SDK for interacting with the `generativelanguage.googleapis.com` API [google-generativeai](#) reads like it was written by a Java developer who thought they knew everything about OOP, spent 30 minutes trying to learn Python, gave up and decided to build the library to prove how horrible Python is. It also doesn't use `httpx` for HTTP requests, and tries to implement tool calling itself, but doesn't use Pydantic or equivalent for validation.

We therefore implement support for the API directly.

Despite these shortcomings, the Gemini model is actually quite powerful and very fast.

Setup

For details on how to set up authentication with this model, see [model configuration for Gemini](#).

GeminiModelName module-attribute

```
GeminiModelName = Literal[
    "gemini-1.5-flash",
    "gemini-1.5-flash-8b",
    "gemini-1.5-pro",
    "gemini-1.0-pro",
]
```

Named Gemini models.

See [the Gemini API docs](#) for a full list.

GeminiModel dataclass

Bases: [Model](#)

A model that uses Gemini via `generativelanguage.googleapis.com` API.

This is implemented from scratch rather than using a dedicated SDK, good API documentation is available [here](#).

Apart from `__init__`, all methods are private or match those of the base class.

```
99 Source code in pydantic_ai_slim/pydantic_ai/models/gemini.py

47 @dataclass(init=False)
48 class GeminiModel(Model):
49     """A model that uses Gemini via 'generativelanguage.googleapis.com' API.
50
51     This is implemented from scratch rather than using a dedicated SDK, good API documentation is
52     available [here](https://ai.google.dev/api).
53
54     Apart from '.__init__', all methods are private or match those of the base class.
55     """
56
57     model_name: GeminiModelName
58     auth: AuthProtocol
59     http_client: AsyncHTTPClient
60     url: str
61
62     def __init__(
63         self,
64         model_name: GeminiModelName,
65         *,
66         api_key: str | None = None,
67         http_client: AsyncHTTPClient | None = None,
68         url_template: str = 'https://generativelanguage.googleapis.com/v1beta/models/{model}:',
69     ):
70         """Initialize a Gemini model.
71
72         Args:
73             model_name: The name of the model to use.
74             api_key: The API key to use for authentication, if not provided, the 'GEMINI_API_KEY' environment variable
75                     will be used if available.
76             http_client: An existing 'httpx.AsyncClient' to use for making HTTP requests.
77             url_template: The URL template to use for making requests, you shouldn't need to change this,
78                         docs [here](https://ai.google.dev/gemini-api/docs/quickstart?lang=rest#make-first-request),
79                         'model' is substituted with the model name, and 'function' is added to the end of the URL.
80         """
81         self.model_name = model_name
82         if api_key is None:
83             if env_api_key := os.getenv('GEMINI_API_KEY'):
84                 api_key = env_api_key
85             else:
86                 raise exceptions.UserError('API key must be provided or set in the GEMINI_API_KEY environment variable')
87         self.auth = ApiKeyAuth(api_key)
88         self.http_client = http_client or cached_async_http_client()
89         self.url = url_template.format(model=model_name)
90
91     async def agent_model(
92         self,
93         *,
94         function_tools: list[ToolDefinition],
95         allow_text_result: bool,
96         result_tools: list[ToolDefinition],
97     ) -> GeminiAgentModel:
98         return GeminiAgentModel(
99             http_client=self.http_client,
100             model_name=self.model_name,
101             auth=self.auth,
102             url=self.url,
103             function_tools=function_tools,
104             allow_text_result=allow_text_result,
105             result_tools=result_tools,
106         )
107
108     def name(self) -> str:
109         return self.model_name
```

`__init__`

```
__init__(
    model_name: GeminiModelName,
    *,
    api_key: str | None = None,
    http_client: AsyncClient | None = None,
    url_template: str = "https://generativelanguage.googleapis.com/v1beta/models/{model}:"
)
```

Initialize a Gemini model.

Parameters:

Name	Type	Description	Default
model_name	GeminiModelName	The name of the model to use.	required
api_key	str None	The API key to use for authentication, if not provided, the GEMINI_API_KEY environment variable will be used if available.	None
http_client	AsyncClient None	An existing httpx.AsyncClient to use for making HTTP requests.	None
url_template	str	The URL template to use for making requests, you shouldn't need to change this, docs here, model is substituted with the model name, and function is added to the end of the URL.	'https://generativelanguage.googleapis.com/v1beta/models/{model}:'

Source code in pydantic_ai_slim/pydantic_ai/models/gemini.py

```
62 def __init__(
63     self,
64     model_name: GeminiModelName,
65     *,
66     api_key: str | None = None,
67     http_client: AsyncHTTPClient | None = None,
68     url_template: str = 'https://generativelanguage.googleapis.com/v1beta/models/{model}:',
69 ):
70     """Initialize a Gemini model.
71
72     Args:
73         model_name: The name of the model to use.
74         api_key: The API key to use for authentication, if not provided, the 'GEMINI_API_KEY' environment variable
75             will be used if available.
76         http_client: An existing 'httpx.AsyncClient' to use for making HTTP requests.
77         url_template: The URL template to use for making requests, you shouldn't need to change this,
78             docs [here](https://ai.google.dev/gemini-api/docs/quickstart?lang=rest#make-first-request),
79             'model' is substituted with the model name, and 'function' is added to the end of the URL.
80
81     """
82     self.model_name = model_name
83     if api_key is None:
84         if env_api_key := os.getenv('GEMINI_API_KEY'):
85             api_key = env_api_key
86         else:
87             raise exceptions.UserError('API key must be provided or set in the GEMINI_API_KEY environment variable')
88     self.auth = ApiKeyAuth(api_key)
89     self.http_client = http_client or cached_async_http_client()
90     self.url = url_template.format(model=model_name)
```

AuthProtocol

Bases: Protocol

Abstract definition for Gemini authentication.

Source code in pydantic_ai_slim/pydantic_ai/models/gemini.py

```
112 class AuthProtocol(Protocol):
113     """Abstract definition for Gemini authentication."""
114
115     async def headers(self) -> dict[str, str]: ...
```

ApiKeyAuth dataclass

Authentication using an API key for the X-Goog-API-Key header.

Source code in pydantic_ai_slim/pydantic_ai/models/gemini.py

```
118 @dataclass
119 class ApiKeyAuth:
120     """Authentication using an API key for the 'X-Goog-API-Key' header."""
121
122     api_key: str
123
124     async def headers(self) -> dict[str, str]:
125         # https://cloud.google.com/docs/authentication/api-keys-use#using-with-rest
126         return {'X-Goog-API-Key': self.api_key}
```

GeminiAgentModel dataclass

Bases: AgentModel

Implementation of AgentModel for Gemini models.

```

129 @dataclass(init=False)
130 class GeminiAgentModel(AgentModel):
131     """Implementation of 'AgentModel' for Gemini models."""
132
133     http_client: AsyncHTTPClient
134     model_name: GeminiModelName
135     auth: AuthProtocol
136     tools: _GeminiTools | None
137     tool_config: _GeminiToolConfig | None
138     url: str
139
140     def __init__(
141         self,
142         http_client: AsyncHTTPClient,
143         model_name: GeminiModelName,
144         auth: AuthProtocol,
145         url: str,
146         function_tools: list[ToolDefinition],
147         allow_text_result: bool,
148         result_tools: list[ToolDefinition],
149     ):
150         check_allow_model_requests()
151         tools = [_function_from_abstract_tool(t) for t in function_tools]
152         if result_tools:
153             tools += [_function_from_abstract_tool(t) for t in result_tools]
154
155         if allow_text_result:
156             tool_config = None
157         else:
158             tool_config = _tool_config([t['name'] for t in tools])
159
160         self.http_client = http_client
161         self.model_name = model_name
162         self.auth = auth
163         self.tools = _GeminiTools(function_declarations=tools) if tools else None
164         self.tool_config = tool_config
165         self.url = url
166
167     async def request(self, messages: list[Message]) -> tuple[ModelAnyResponse, result.Cost]:
168         async with self._make_request(messages, False) as http_response:
169             response = _gemini_response_ta.validate_json(await http_response.aread())
170             return self._process_response(response), _metadata_as_cost(response)
171
172     @asynctxcontextmanager
173     async def request_stream(self, messages: list[Message]) -> AsyncIterator[EitherStreamedResponse]:
174         async with self._make_request(messages, True) as http_response:
175             yield await self._process_streamed_response(http_response)
176
177     @asynctxcontextmanager
178     async def _make_request(self, messages: list[Message], streamed: bool) -> AsyncIterator[HTTPResponse]:
179         contents: list[_GeminiContent] = []
180         sys_prompt_parts: list[_GeminiTextPart] = []
181         for m in messages:
182             either_content = self._message_to_gemini(m)
183             if left := either_content.left:
184                 sys_prompt_parts.append(left.value)
185             else:
186                 contents.append(either_content.right)
187
188         request_data = _GeminiRequest(contents=contents)
189         if sys_prompt_parts:
190             request_data['system_instruction'] = _GeminiTextContent(role='user', parts=sys_prompt_parts)
191         if self.tools is not None:
192             request_data['tools'] = self.tools
193         if self.tool_config is not None:
194             request_data['tool_config'] = self.tool_config
195
196         url = self.url + ('streamGenerateContent' if streamed else 'generateContent')
197
198         headers = {
199             'Content-Type': 'application/json',
200             'User-Agent': get_user_agent(),
201             **await self.auth.headers(),
202         }
203
204         request_json = _gemini_request_ta.dump_json(request_data, by_alias=True)
205
206         async with self.http_client.stream('POST', url, content=request_json, headers=headers) as r:
207             if r.status_code != 200:
208                 await r.aread()
209                 raise exceptions.UnexpectedModelBehavior(f'Unexpected response from gemini {r.status_code}', r.text)
210             yield r
211
212     @staticmethod
213     def _process_response(response: _GeminiResponse) -> ModelAnyResponse:
214         either = _extract_response_parts(response)
215         if left := either.left:
216             return _structured_response_from_parts(left.value)
217         else:
218             return ModelTextResponse(content=''.join(part['text'] for part in either.right))
219
220     @staticmethod
221     async def _process_streamed_response(http_response: HTTPResponse) -> EitherStreamedResponse:
222         """Process a streamed response, and prepare a streaming response to return."""
223         aiter_bytes = http_response.aiter_bytes()
224         start_response: _GeminiResponse | None = None
225         content = bytearray()
226
227         async for chunk in aiter_bytes:
228             content.extend(chunk)
229             responses = _gemini_streamed_response_ta.validate_json(
230                 content,
231                 experimental_allow_partial='trailing-strings',
232             )
233             if responses:
234                 last = responses[-1]
235                 if last['candidates'] and last['candidates'][0]['content']['parts']:
236                     start_response = last
237                     break
238
239         if start_response is None:
240             raise UnexpectedModelBehavior('Streamed response ended without content or tool calls')
241
242         if _extract_response_parts(start_response).is_left():
243             return GeminiStreamStructuredResponse(_content=content, _stream=aiter_bytes)
244         else:
245             return GeminiStreamTextResponse(_json_content=content, _stream=aiter_bytes)
246
247     @staticmethod
248     def _message_to_gemini(m: Message) -> _utils.Either[_GeminiTextPart, _GeminiContent]:
249         """Convert a message to a _GeminiTextPart for "system_instructions" or _GeminiContent for "contents"."""
250         if m.role == 'system':
251             # SystemPrompt ->
252             return _utils.Either(left=_GeminiTextPart(text=m.content))
253         elif m.role == 'user':
254             # UserPrompt ->
255             return _utils.Either(right=_content_user_text(m.content))

```

```

256         elif m.role == 'tool-return':
257             # ToolReturn ->
258             return _utils.Either(right=_content_function_return(m))
259         elif m.role == 'retry-prompt':
260             # RetryPrompt ->
261             return _utils.Either(right=_content_function_retry(m))
262         elif m.role == 'model-text-response':
263             # ModelTextResponse ->
264             return _utils.Either(right=_content_model_text(m.content))
265         elif m.role == 'model-structured-response':
266             # ModelStructuredResponse ->
267             return _utils.Either(right=_content_function_call(m))
268         else:
269             assert_never(m)

```

GeminiStreamTextResponse `dataclass`

Bases: `StreamTextResponse`

Implementation of `StreamTextResponse` for the Gemini model.

```

99 Source code in pydantic_ai_slim/pydantic_ai/models/gemini.py
272 @dataclass
273 class GeminiStreamTextResponse(StreamTextResponse):
274     """Implementation of 'StreamTextResponse' for the Gemini model."""
275
276     _json_content: bytearray
277     _stream: AsyncIterator[bytes]
278     _position: int = 0
279     _timestamp: datetime = field(default_factory=_utils.now_utc, init=False)
280     _cost: result.Cost = field(default_factory=result.Cost, init=False)
281
282     async def __anext__(self) -> None:
283         chunk = await self._stream.__anext__()
284         self._json_content.extend(chunk)
285
286     def get(self, *, final: bool = False) -> Iterable[str]:
287         if final:
288             all_items = pydantic_core.from_json(self._json_content)
289             new_items = all_items[self._position :]
290             self._position = len(all_items)
291             new_responses = _gemini_streamed_response_ta.validate_python(new_items)
292         else:
293             all_items = pydantic_core.from_json(self._json_content, allow_partial=True)
294             new_items = all_items[self._position : -1]
295             self._position = len(all_items) - 1
296             new_responses = _gemini_streamed_response_ta.validate_python(
297                 new_items, experimental_allow_partial='trailing-strings'
298             )
299         for r in new_responses:
300             self._cost += _metadata_as_cost(r)
301             parts = r['candidates'][0]['content']['parts']
302             if _all_text_parts(parts):
303                 for part in parts:
304                     yield part['text']
305             else:
306                 raise UnexpectedModelBehavior(
307                     'Streamed response with unexpected content, expected all parts to be text'
308                 )
309
310     def cost(self) -> result.Cost:
311         return self._cost
312
313     def timestamp(self) -> datetime:
314         return self._timestamp

```

GeminiStreamStructuredResponse `dataclass`

Bases: `StreamStructuredResponse`

Implementation of `StreamStructuredResponse` for the Gemini model.

```

317 @dataclass
318 class GeminiStreamStructuredResponse(StreamStructuredResponse):
319     """Implementation of 'StreamStructuredResponse' for the Gemini model."""
320
321     _content: bytearray
322     _stream: AsyncIterator[bytes]
323     _timestamp: datetime = field(default_factory=_utils.now_utc, init=False)
324     _cost: result.Cost = field(default_factory=result.Cost, init=False)
325
326     async def __anext__(self) -> None:
327         chunk = await self._stream.__anext__()
328         self._content.extend(chunk)
329
330     def get(self, *, final: bool = False) -> ModelStructuredResponse:
331         """Get the 'ModelStructuredResponse' at this point.
332
333         NOTE: It's not clear how the stream of responses should be combined because Gemini seems to always
334         reply with a single response, when returning a structured data.
335
336         I'm therefore assuming that each part contains a complete tool call, and not trying to combine data from
337         separate parts.
338         """
339         responses = _gemini_streamed_response_ta.validate_json(
340             self._content,
341             experimental_allow_partial='off' if final else 'trailing-strings',
342         )
343         combined_parts: list[_GeminiFunctionCallPart] = []
344         self._cost = result.Cost()
345         for r in responses:
346             self._cost += _metadata_as_cost(r)
347             candidate = r['candidates'][0]
348             parts = candidate['content']['parts']
349             if _all_function_call_parts(parts):
350                 combined_parts.extend(parts)
351             elif not candidate.get('finish_reason'):
352                 # you can get an empty text part along with the finish_reason, so we ignore that case
353                 raise UnexpectedModelBehavior(
354                     'Streamed response with unexpected content, expected all parts to be function calls'
355                 )
356         return _structured_response_from_parts(combined_parts, timestamp=self._timestamp)
357
358     def cost(self) -> result.Cost:
359         return self._cost
360
361     def timestamp(self) -> datetime:
362         return self._timestamp

```

get

```

get(*, final: bool = False) -> ModelStructuredResponse

```

Get the `ModelStructuredResponse` at this point.

NOTE: It's not clear how the stream of responses should be combined because Gemini seems to always reply with a single response, when returning a structured data.

I'm therefore assuming that each part contains a complete tool call, and not trying to combine data from separate parts.

```

330     def get(self, *, final: bool = False) -> ModelStructuredResponse:
331         """Get the 'ModelStructuredResponse' at this point.
332
333         NOTE: It's not clear how the stream of responses should be combined because Gemini seems to always
334         reply with a single response, when returning a structured data.
335
336         I'm therefore assuming that each part contains a complete tool call, and not trying to combine data from
337         separate parts.
338         """
339         responses = _gemini_streamed_response_ta.validate_json(
340             self._content,
341             experimental_allow_partial='off' if final else 'trailing-strings',
342         )
343         combined_parts: list[_GeminiFunctionCallPart] = []
344         self._cost = result.Cost()
345         for r in responses:
346             self._cost += _metadata_as_cost(r)
347             candidate = r['candidates'][0]
348             parts = candidate['content']['parts']
349             if _all_function_call_parts(parts):
350                 combined_parts.extend(parts)
351             elif not candidate.get('finish_reason'):
352                 # you can get an empty text part along with the finish_reason, so we ignore that case
353                 raise UnexpectedModelBehavior(
354                     'Streamed response with unexpected content, expected all parts to be function calls'
355                 )
356         return _structured_response_from_parts(combined_parts, timestamp=self._timestamp)

```