

pydantic_ai.models.openai

Setup

For details on how to set up authentication with this model, see [model configuration for OpenAI](#).

OpenAIModelName module-attribute

```
OpenAIModelName = Union[ChatModel, str]
```

Using this more broad type for the model name instead of the ChatModel definition allows this model to be used more easily with other model types (ie, Ollama)

OpenAIModel dataclass

Bases: [Model](#)

A model that uses the OpenAI API.

Internally, this uses the [OpenAI Python client](#) to interact with the API.

Apart from `__init__`, all methods are private or match those of the base class.

```

99 Source code in pydantic_ai_slim/pydantic_ai/models/openai.py
53 @dataclass(init=False)
54 class OpenAIModel(Model):
55     """A model that uses the OpenAI API.
56
57     Internally, this uses the [OpenAI Python client](https://github.com/openai/openai-python) to interact with the API.
58
59     Apart from `__init__`, all methods are private or match those of the base class.
60     """
61
62     model_name: OpenAIModelName
63     client: AsyncOpenAI = field(repr=False)
64
65     def __init__(
66         self,
67         model_name: OpenAIModelName,
68         *,
69         api_key: str | None = None,
70         openai_client: AsyncOpenAI | None = None,
71         http_client: AsyncHTTPClient | None = None,
72     ):
73         """Initialize an OpenAI model.
74
75         Args:
76             model_name: The name of the OpenAI model to use. List of model names available
77                 [here](https://github.com/openai/openai-python/blob/v1.54.3/src/openai/types/chat_model.py#L7)
78                 (Unfortunately, despite being asked to do so, OpenAI do not provide '.inv' files for their API).
79             api_key: The API key to use for authentication, if not provided, the 'OPENAI_API_KEY' environment variable
80                 will be used if available.
81             openai_client: An existing
82                 ['AsyncOpenAI'](https://github.com/openai/openai-python?tab=readme-ov-file#async-usage)
83                 client to use, if provided, 'api_key' and 'http_client' must be 'None'.
84             http_client: An existing 'httpx.AsyncClient' to use for making HTTP requests.
85
86         """
87         self.model_name = OpenAIModelName = model_name
88         if openai_client is not None:
89             assert http_client is None, 'Cannot provide both `openai_client` and `http_client`'
90             assert api_key is None, 'Cannot provide both `openai_client` and `api_key`'
91             self.client = openai_client
92         elif http_client is not None:
93             self.client = AsyncOpenAI(api_key=api_key, http_client=http_client)
94         else:
95             self.client = AsyncOpenAI(api_key=api_key, http_client=cached_async_http_client())
96
97     async def agent_model(
98         self,
99         *,
100         function_tools: list[ToolDefinition],
101         allow_text_result: bool,
102         result_tools: list[ToolDefinition],
103     ) -> AgentModel:
104         check_allow_model_requests()
105         tools = [self._map_tool_definition(r) for r in function_tools]
106         if result_tools:
107             tools += [self._map_tool_definition(r) for r in result_tools]
108         return OpenAIAgentModel(
109             self.client,
110             self.model_name,
111             allow_text_result,
112             tools,
113         )
114
115     def name(self) -> str:
116         return f'openai:{self.model_name}'
117
118     @staticmethod
119     def _map_tool_definition(f: ToolDefinition) -> chat.ChatCompletionToolParam:
120         return {
121             'type': 'function',
122             'function': {
123                 'name': f.name,
124                 'description': f.description,
125                 'parameters': f.parameters_json_schema,
126             },
127         }

```

__init__

```

__init__(
    self,
    model_name: OpenAIModelName,
    *,
    api_key: str | None = None,
    openai_client: AsyncOpenAI | None = None,
    http_client: AsyncClient | None = None
)

```

Parameters:

Name	Type	Description	Default
<code>model_name</code>	<code>OpenAIModelName</code>	The name of the OpenAI model to use. List of model names available here (Unfortunately, despite being ask to do so, OpenAI do not provide <code>.inv</code> files for their API).	<i>required</i>
<code>api_key</code>	<code>str</code> <code>None</code>	The API key to use for authentication, if not provided, the <code>OPENAI_API_KEY</code> environment variable will be used if available.	<code>None</code>
<code>openai_client</code>	<code>AsyncOpenAI</code> <code>None</code>	An existing <code>AsyncOpenAI</code> client to use, if provided, <code>api_key</code> and <code>http_client</code> must be <code>None</code> .	<code>None</code>
<code>http_client</code>	<code>AsyncClient</code> <code>None</code>	An existing <code>httpx.AsyncClient</code> to use for making HTTP requests.	<code>None</code>

99 Source code in `pydantic_ai_slim/pydantic_ai/models/openai.py`

```
65 def __init__(
66     self,
67     model_name: OpenAIModelName,
68     *,
69     api_key: str | None = None,
70     openai_client: AsyncOpenAI | None = None,
71     http_client: AsyncHTTPClient | None = None,
72 ):
73     """Initialize an OpenAI model.
74
75     Args:
76         model_name: The name of the OpenAI model to use. List of model names available
77             [here](https://github.com/openai/openai-python/blob/v1.54.3/src/openai/types/chat_model.py#L7)
78             (Unfortunately, despite being ask to do so, OpenAI do not provide .inv files for their API).
79         api_key: The API key to use for authentication, if not provided, the 'OPENAI_API_KEY' environment variable
80             will be used if available.
81         openai_client: An existing
82             ['AsyncOpenAI'](https://github.com/openai/openai-python?tab=readme-ov-file#async-usage)
83             client to use, if provided, 'api_key' and 'http_client' must be 'None'.
84         http_client: An existing 'httpx.AsyncClient' to use for making HTTP requests.
85     """
86     self.model_name = OpenAIModelName = model_name
87     if openai_client is not None:
88         assert http_client is None, 'Cannot provide both 'openai_client' and 'http_client''
89         assert api_key is None, 'Cannot provide both 'openai_client' and 'api_key''
90         self.client = openai_client
91     elif http_client is not None:
92         self.client = AsyncOpenAI(api_key=api_key, http_client=http_client)
93     else:
94         self.client = AsyncOpenAI(api_key=api_key, http_client=cached_async_http_client())
```

OpenAI AgentModel `dataclass`

Bases: `AgentModel`

Implementation of `AgentModel` for OpenAI models.

```

129 @dataclass
130 class OpenAIModel(AgentModel):
131     """Implementation of 'AgentModel' for OpenAI models."""
132
133     client: AsyncOpenAI
134     model_name: OpenAIModelName
135     allow_text_result: bool
136     tools: list[chat.ChatCompletionToolParam]
137
138     async def request(self, messages: list[Message]) -> tuple[ModelAnyResponse, result.Cost]:
139         response = await self._completions_create(messages, False)
140         return self._process_response(response), _map_cost(response)
141
142     @asynctxcontextmanager
143     async def request_stream(self, messages: list[Message]) -> AsyncIterator[EitherStreamedResponse]:
144         response = await self._completions_create(messages, True)
145         async with response:
146             yield await self._process_streamed_response(response)
147
148     @overload
149     async def _completions_create(
150         self, messages: list[Message], stream: Literal[True]
151     ) -> AsyncStream[ChatCompletionChunk]:
152         pass
153
154     @overload
155     async def _completions_create(self, messages: list[Message], stream: Literal[False]) -> chat.ChatCompletion:
156         pass
157
158     async def _completions_create(
159         self, messages: list[Message], stream: bool
160     ) -> chat.ChatCompletion | AsyncStream[ChatCompletionChunk]:
161         # standalone function to make it easier to override
162         if not self.tools:
163             tool_choice: Literal['none', 'required', 'auto'] | None = None
164         elif not self.allow_text_result:
165             tool_choice = 'required'
166         else:
167             tool_choice = 'auto'
168
169         openai_messages = [self._map_message(m) for m in messages]
170         return await self.client.chat.completions.create(
171             model=self.model_name,
172             messages=openai_messages,
173             n=1,
174             parallel_tool_calls=True if self.tools else NOT_GIVEN,
175             tools=self.tools or NOT_GIVEN,
176             tool_choice=tool_choice or NOT_GIVEN,
177             stream=stream,
178             stream_options={'include_usage': True} if stream else NOT_GIVEN,
179         )
180
181     @staticmethod
182     def _process_response(response: chat.ChatCompletion) -> ModelAnyResponse:
183         """Process a non-streamed response, and prepare a message to return."""
184         timestamp = datetime.fromtimestamp(response.created, tz=timezone.utc)
185         choice = response.choices[0]
186         if choice.message.tool_calls is not None:
187             return ModelStructuredResponse(
188                 [ToolCall.from_json(c.function.name, c.function.arguments, c.id) for c in choice.message.tool_calls],
189                 timestamp=timestamp,
190             )
191         else:
192             assert choice.message.content is not None, choice
193             return ModelTextResponse(choice.message.content, timestamp=timestamp)
194
195     @staticmethod
196     async def _process_streamed_response(response: AsyncStream[ChatCompletionChunk]) -> EitherStreamedResponse:
197         """Process a streamed response, and prepare a streaming response to return."""
198         timestamp: datetime | None = None
199         start_cost = Cost()
200         # the first chunk may contain enough information so we iterate until we get either 'tool_calls' or 'content'
201         while True:
202             try:
203                 chunk = await response._anext__()
204             except StopAsyncIteration as e:
205                 raise UnexpectedModelBehavior('Streamed response ended without content or tool calls') from e
206
207             timestamp = timestamp or datetime.fromtimestamp(chunk.created, tz=timezone.utc)
208             start_cost += _map_cost(chunk)
209
210             if chunk.choices:
211                 delta = chunk.choices[0].delta
212
213                 if delta.content is not None:
214                     return OpenAIStreamTextResponse(delta.content, response, timestamp, start_cost)
215                 elif delta.tool_calls is not None:
216                     return OpenAIStreamStructuredResponse(
217                         response,
218                         {c.index: c for c in delta.tool_calls},
219                         timestamp,
220                         start_cost,
221                     )
222                 # else continue until we get either delta.content or delta.tool_calls
223
224     @staticmethod
225     def _map_message(message: Message) -> chat.ChatCompletionMessageParam:
226         """Just maps a 'pydantic_ai.Message' to a 'openai.types.ChatCompletionMessageParam'."""
227         if message.role == 'system':
228             # SystemPrompt ->
229             return chat.ChatCompletionSystemMessageParam(role='system', content=message.content)
230         elif message.role == 'user':
231             # UserPrompt ->
232             return chat.ChatCompletionUserMessageParam(role='user', content=message.content)
233         elif message.role == 'tool-return':
234             # ToolReturn ->
235             return chat.ChatCompletionToolMessageParam(
236                 role='tool',
237                 tool_call_id=_guard_tool_call_id(message),
238                 content=message.model_response_str(),
239             )
240         elif message.role == 'retry-prompt':
241             # RetryPrompt ->
242             if message.tool_name is None:
243                 return chat.ChatCompletionUserMessageParam(role='user', content=message.model_response())
244             else:
245                 return chat.ChatCompletionToolMessageParam(
246                     role='tool',
247                     tool_call_id=_guard_tool_call_id(message),
248                     content=message.model_response(),
249                 )
250         elif message.role == 'model-text-response':
251             # ModelTextResponse ->
252             return chat.ChatCompletionAssistantMessageParam(role='assistant', content=message.content)
253         elif message.role == 'model-structured-response':
254             assert (
255                 message.role == 'model-structured-response'

```

```

256         ), f'Expected role to be "llm-tool-calls", got {message.role}'
257         # ModelStructuredResponse ->
258         return chat.ChatCompletionAssistantMessageParam(
259             role='assistant',
260             tool_calls=[_map_tool_call(t) for t in message.calls],
261         )
262     else:
263         assert_never(message)

```

OpenAIStreamTextResponse `dataclass`

Bases: `StreamTextResponse`

Implementation of `StreamTextResponse` for OpenAI models.

99 Source code in `pydantic_ai_slim/pydantic_ai/models/openai.py`

```

266 @dataclass
267 class OpenAIStreamTextResponse(StreamTextResponse):
268     """Implementation of 'StreamTextResponse' for OpenAI models."""
269
270     _first: str | None
271     _response: AsyncStream[ChatCompletionChunk]
272     _timestamp: datetime
273     _cost: result.Cost
274     _buffer: list[str] = field(default_factory=list, init=False)
275
276     async def __anext__(self) -> None:
277         if self._first is not None:
278             self._buffer.append(self._first)
279             self._first = None
280             return None
281
282         chunk = await self._response.__anext__()
283         self._cost += _map_cost(chunk)
284         try:
285             choice = chunk.choices[0]
286         except IndexError:
287             raise StopAsyncIteration()
288
289         # we don't raise StopAsyncIteration on the last chunk because usage comes after this
290         if choice.finish_reason is None:
291             assert choice.delta.content is not None, f'Expected delta with content, invalid chunk: {chunk!r}'
292             if choice.delta.content is not None:
293                 self._buffer.append(choice.delta.content)
294
295     def get(self, *, final: bool = False) -> Iterable[str]:
296         yield from self._buffer
297         self._buffer.clear()
298
299     def cost(self) -> Cost:
300         return self._cost
301
302     def timestamp(self) -> datetime:
303         return self._timestamp

```

OpenAIStreamStructuredResponse `dataclass`

Bases: `StreamStructuredResponse`

Implementation of `StreamStructuredResponse` for OpenAI models.

99 Source code in `pydantic_ai_slim/pydantic_ai/models/openai.py`

```

306 @dataclass
307 class OpenAIStreamStructuredResponse(StreamStructuredResponse):
308     """Implementation of 'StreamStructuredResponse' for OpenAI models."""
309
310     _response: AsyncStream[ChatCompletionChunk]
311     _delta_tool_calls: dict[int, ChoiceDeltaToolCall]
312     _timestamp: datetime
313     _cost: result.Cost
314
315     async def __anext__(self) -> None:
316         chunk = await self._response.__anext__()
317         self._cost += _map_cost(chunk)
318         try:
319             choice = chunk.choices[0]
320         except IndexError:
321             raise StopAsyncIteration()
322
323         if choice.finish_reason is not None:
324             raise StopAsyncIteration()
325
326         assert choice.delta.content is None, f'Expected tool calls, got content instead, invalid chunk: {chunk!r}'
327
328         for new in choice.delta.tool_calls or []:
329             if current := self._delta_tool_calls.get(new.index):
330                 if current.function is None:
331                     current.function = new.function
332                 elif new.function is not None:
333                     current.function.name = _utils.add_optional(current.function.name, new.function.name)
334                     current.function.arguments = _utils.add_optional(current.function.arguments, new.function.arguments)
335             else:
336                 self._delta_tool_calls[new.index] = new
337
338     def get(self, *, final: bool = False) -> ModelStructuredResponse:
339         calls: list[ToolCall] = []
340         for c in self._delta_tool_calls.values():
341             if f := c.function:
342                 if f.name is not None and f.arguments is not None:
343                     calls.append(ToolCall.from_json(f.name, f.arguments, c.id))
344
345         return ModelStructuredResponse(calls, timestamp=self._timestamp)
346
347     def cost(self) -> Cost:
348         return self._cost
349
350     def timestamp(self) -> datetime:
351         return self._timestamp

```