

## pydantic\_ai.result

**ResultData** module-attribute

```
ResultData = TypeVar('ResultData')
```

Type variable for the result data of a run.

**RunResult** dataclass

Bases: `_BaseRunResult[ResultData]`

Result of a non-streamed run.

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```

102 @dataclass
103 class RunResult(_BaseRunResult[ResultData]):
104     """Result of a non-streamed run."""
105
106     data: ResultData
107     """Data from the final response in the run."""
108     _cost: Cost
109
110     def cost(self) -> Cost:
111         """Return the cost of the whole run."""
112         return self._cost

```

**all\_messages**

```
all_messages() -> list[Message]
```

Return the history of messages.

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```

77 def all_messages(self) -> list[messages.Message]:
78     """Return the history of messages."""
79     # this is a method to be consistent with the other methods
80     return self._all_messages

```

**all\_messages\_json**

```
all_messages_json() -> bytes
```

Return all messages from `all_messages` as JSON bytes.

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```

82 def all_messages_json(self) -> bytes:
83     """Return all messages from ['all_messages'] as JSON bytes."""
84     return messages.MessagesTypeAdapter.dump_json(self.all_messages())

```

**new\_messages**

```
new_messages() -> list[Message]
```

Return new messages associated with this run.

System prompts and any messages from older runs are excluded.

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```

86 def new_messages(self) -> list[messages.Message]:
87     """Return new messages associated with this run.
88
89     System prompts and any messages from older runs are excluded.
90     """
91     return self.all_messages()[self._new_message_index :]

```

**new\_messages\_json**

```
new_messages_json() -> bytes
```

Return new messages from `new_messages` as JSON bytes.

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```

93 def new_messages_json(self) -> bytes:
94     """Return new messages from ['new_messages'] as JSON bytes."""
95     return messages.MessagesTypeAdapter.dump_json(self.new_messages())

```

**data** instance-attribute

```
data: ResultData
```

Data from the final response in the run.

cost

```
cost() -> Cost
```



Return the cost of the whole run.

Source code in pydantic\_ai\_slim/pydantic\_ai/result.py



```
110 def cost(self) -> Cost:
111     """Return the cost of the whole run."""
112     return self._cost
```



StreamedRunResult `dataclass`

Bases: `_BaseRunResult[ResultData]`, `Generic[AgentDeps, ResultData]`

Result of a streamed run that returns structured data via a tool call.

```

115 @dataclass
116 class StreamedRunResult(BaseRunResult[ResultData], Generic[AgentDeps, ResultData]):
117     """Result of a streamed run that returns structured data via a tool call."""
118
119     cost_so_far: Cost
120     """Cost of the run up until the last request."""
121     _stream_response: models.EitherStreamedResponse
122     _result_schema: _result.ResultSchema[ResultData] | None
123     _deps: AgentDeps
124     _result_validators: list[_result.ResultValidator[AgentDeps, ResultData]]
125     _on_complete: Callable[[list[messages.Message]], None]
126     is_complete: bool = field(default=False, init=False)
127     """Whether the stream has all been received.
128
129     This is set to 'True' when one of
130     ['stream'] [pydantic_ai.result.StreamedRunResult.stream],
131     ['stream_text'] [pydantic_ai.result.StreamedRunResult.stream_text],
132     ['stream_structured'] [pydantic_ai.result.StreamedRunResult.stream_structured] or
133     ['get_data'] [pydantic_ai.result.StreamedRunResult.get_data] completes.
134     """
135
136     async def stream(self, *, debounce_by: float | None = 0.1) -> AsyncIterator[ResultData]:
137         """Stream the response as an async iterable.
138
139         The pydantic validator for structured data will be called in
140         [partial mode] (https://docs.pydantic.dev/dev/concepts/experimental/#partial-validation)
141         on each iteration.
142
143         Args:
144             debounce_by: by how much (if at all) to debounce/group the response chunks by. 'None' means no debouncing.
145                 Debouncing is particularly important for long structured responses to reduce the overhead of
146                 performing validation as each token is received.
147
148         Returns:
149             An async iterable of the response data.
150         """
151         if isinstance(self._stream_response, models.StreamTextResponse):
152             async for text in self.stream_text(debounce_by=debounce_by):
153                 yield cast(ResultData, text)
154         else:
155             async for structured_message, is_last in self.stream_structured(debounce_by=debounce_by):
156                 yield await self.validate_structured_result(structured_message, allow_partial=not is_last)
157
158     async def stream_text(self, *, delta: bool = False, debounce_by: float | None = 0.1) -> AsyncIterator[str]:
159         """Stream the text result as an async iterable.
160
161         !!! note
162             This method will fail if the response is structured,
163             e.g. if ['is_structured'] [pydantic_ai.result.StreamedRunResult.is_structured] returns 'True'.
164
165         !!! note
166             Result validators will NOT be called on the text result if 'delta=True'.
167
168         Args:
169             delta: if 'True', yield each chunk of text as it is received, if 'False' (default), yield the full text
170                 up to the current point.
171             debounce_by: by how much (if at all) to debounce/group the response chunks by. 'None' means no debouncing.
172                 Debouncing is particularly important for long structured responses to reduce the overhead of
173                 performing validation as each token is received.
174         """
175         with _logfire.span('response stream text') as lf_span:
176             if isinstance(self._stream_response, models.StreamStructuredResponse):
177                 raise exceptions.UserError('stream_text() can only be used with text responses')
178             if delta:
179                 async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
180                     async for _ in group_iter:
181                         yield ''.join(self._stream_response.get())
182                     final_delta = ''.join(self._stream_response.get(final=True))
183                     if final_delta:
184                         yield final_delta
185             else:
186                 # a quick benchmark shows it's faster to build up a string with concat when we're
187                 # yielding at each step
188                 chunks: list[str] = []
189                 combined = ''
190                 async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
191                     async for _ in group_iter:
192                         new = False
193                         for chunk in self._stream_response.get():
194                             chunks.append(chunk)
195                             new = True
196                         if new:
197                             combined = await self._validate_text_result(''.join(chunks))
198                             yield combined
199
200                 new = False
201                 for chunk in self._stream_response.get(final=True):
202                     chunks.append(chunk)
203                     new = True
204                 if new:
205                     combined = await self._validate_text_result(''.join(chunks))
206                     yield combined
207                 lf_span.set_attribute('combined_text', combined)
208                 self._marked_completed(text=combined)
209
210     async def stream_structured(
211         self, *, debounce_by: float | None = 0.1
212     ) -> AsyncIterator[tuple[messages.ModelStructuredResponse, bool]]:
213         """Stream the response as an async iterable of Structured LLM Messages.
214
215         !!! note
216             This method will fail if the response is text,
217             e.g. if ['is_structured'] [pydantic_ai.result.StreamedRunResult.is_structured] returns 'False'.
218
219         Args:
220             debounce_by: by how much (if at all) to debounce/group the response chunks by. 'None' means no debouncing.
221                 Debouncing is particularly important for long structured responses to reduce the overhead of
222                 performing validation as each token is received.
223
224         Returns:
225             An async iterable of the structured response message and whether that is the last message.
226         """
227         with _logfire.span('response stream structured') as lf_span:
228             if isinstance(self._stream_response, models.StreamTextResponse):
229                 raise exceptions.UserError('stream_structured() can only be used with structured responses')
230             else:
231                 # we should already have a message at this point, yield that first if it has any content
232                 msg = self._stream_response.get()
233                 if any(call.has_content() for call in msg.calls):
234                     yield msg, False
235                 async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
236                     async for _ in group_iter:
237                         msg = self._stream_response.get()
238                         if any(call.has_content() for call in msg.calls):
239                             yield msg, False
240                 msg = self._stream_response.get(final=True)
241                 yield msg, True

```

```

242         if span.set_attribute('structured_response', msg)
243             self._marked_completed(structured_message=msg)
244
245     async def get_data(self) -> ResultData:
246         """Stream the whole response, validate and return it."""
247         async for _ in self._stream_response:
248             pass
249         if isinstance(self._stream_response, models.StreamTextResponse):
250             text = ''.join(self._stream_response.get(final=True))
251             text = await self._validate_text_result(text)
252             self._marked_completed(text=text)
253             return cast(ResultData, text)
254         else:
255             structured_message = self._stream_response.get(final=True)
256             self._marked_completed(structured_message=structured_message)
257             return await self.validate_structured_result(structured_message)
258
259     @property
260     def is_structured(self) -> bool:
261         """Return whether the stream response contains structured data (as opposed to text)."""
262         return isinstance(self._stream_response, models.StreamStructuredResponse)
263
264     def cost(self) -> Cost:
265         """Return the cost of the whole run.
266
267         !!! note
268             This won't return the full cost until the stream is finished.
269         """
270         return self.cost_so_far + self._stream_response.cost()
271
272     def timestamp(self) -> datetime:
273         """Get the timestamp of the response."""
274         return self._stream_response.timestamp()
275
276     async def validate_structured_result(
277         self, message: messages.ModelStructuredResponse, *, allow_partial: bool = False
278     ) -> ResultData:
279         """Validate a structured result message."""
280         assert self._result_schema is not None, 'Expected _result_schema to not be None'
281         match = self._result_schema.find_tool(message)
282         if match is None:
283             raise exceptions.UnexpectedModelBehavior(
284                 f'Invalid message, unable to find tool: {self._result_schema.tool_names()}'
285             )
286         call, result_tool = match
287         result_data = result_tool.validate(call, allow_partial=allow_partial, wrap_validation_errors=False)
288
289         for validator in self._result_validators:
290             result_data = await validator.validate(result_data, self._deps, 0, call)
291         return result_data
292
293     async def _validate_text_result(self, text: str) -> str:
294         for validator in self._result_validators:
295             text = await validator.validate(
296                 # pyright: ignore[reportAssignmentType]
297                 text,
298                 # pyright: ignore[reportArgumentType]
299                 self._deps,
300                 0,
301                 None,
302             )
303         return text
304
305     def _marked_completed(
306         self, *, text: str | None = None, structured_message: messages.ModelStructuredResponse | None = None
307     ) -> None:
308         self.is_complete = True
309         if text is not None:
310             assert structured_message is None, 'Either text or structured_message should provided, not both'
311             self._all_messages.append(
312                 messages.ModelTextResponse(content=text, timestamp=self._stream_response.timestamp())
313             )
314         else:
315             assert structured_message is not None, 'Either text or structured_message should provided, not both'
316             self._all_messages.append(structured_message)
317         self._on_complete(self._all_messages)

```

## all\_messages

```
all_messages() -> list[Message]
```

Return the history of messages.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```

77 def all_messages(self) -> list[messages.Message]:
78     """Return the history of messages."""
79     # this is a method to be consistent with the other methods
80     return self._all_messages

```

## all\_messages\_json

```
all_messages_json() -> bytes
```

Return all messages from `all_messages` as JSON bytes.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```

82 def all_messages_json(self) -> bytes:
83     """Return all messages from [all_messages][all_messages] as JSON bytes."""
84     return messages.MessagesTypeAdapter.dump_json(self.all_messages())

```

## new\_messages

```
new_messages() -> list[Message]
```

Return new messages associated with this run.

System prompts and any messages from older runs are excluded.

Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
86 def new_messages(self) -> list[messages.Message]:
87     """Return new messages associated with this run.
88
89     System prompts and any messages from older runs are excluded.
90     """
91     return self.all_messages()[self._new_message_index :]
```

new\_messages\_json

```
new_messages_json() -> bytes
```

Return new messages from `new_messages` as JSON bytes.

Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
93 def new_messages_json(self) -> bytes:
94     """Return new messages from ['new_messages'][..new_messages] as JSON bytes."""
95     return messages.MessagesTypeAdapter.dump_json(self.new_messages())
```

cost\_so\_far instance-attribute

```
cost_so_far: Cost
```

Cost of the run up until the last request.

is\_complete class-attribute instance-attribute

```
is_complete: bool = field(default=False, init=False)
```

Whether the stream has all been received.

This is set to `True` when one of `stream`, `stream_text`, `stream_structured` or `get_data` completes.

stream async

```
stream(
    *, debounce_by: float | None = 0.1
) -> AsyncIterator[ResultData]
```

Stream the response as an async iterable.

The pydantic validator for structured data will be called in `partial mode` on each iteration.

Parameters:

| Name        | Type         | Description  | Default |
|-------------|--------------|--|---------|
| debounce_by | float   None | by how much (if at all) to debounce/group the response chunks by. <code>None</code> means no debouncing. Debouncing is particularly important for long structured responses to reduce the overhead of performing validation as each token is received. | 0.1     |

Returns:

| Type                      | Description                             |
|---------------------------|---|
| AsyncIterator[ResultData] | An async iterable of the response data. |

Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
136 async def stream(self, *, debounce_by: float | None = 0.1) -> AsyncIterator[ResultData]:
137     """Stream the response as an async iterable.
138
139     The pydantic validator for structured data will be called in
140     [partial mode](https://docs.pydantic.dev/dev/concepts/experimental/#partial-validation)
141     on each iteration.
142
143     Args:
144         debounce_by: by how much (if at all) to debounce/group the response chunks by. 'None' means no debouncing.
145         Debouncing is particularly important for long structured responses to reduce the overhead of
146         performing validation as each token is received.
147
148     Returns:
149         An async iterable of the response data.
150     """
151     if isinstance(self._stream_response, models.StreamTextResponse):
152         async for text in self.stream_text(debounce_by=debounce_by):
153             yield cast(ResultData, text)
154     else:
155         async for structured_message, is_last in self.stream_structured(debounce_by=debounce_by):
156             yield await self.validate_structured_result(structured_message, allow_partial=not is_last)
```

stream\_text async

```
stream_text(
    *, delta: bool = False, debounce_by: float | None = 0.1
) -> AsyncIterator[str]
```

Stream the text result as an async iterable.

Note

This method will fail if the response is structured, e.g. if `is_structured` returns `True`.

Note

Result validators will NOT be called on the text result if `delta=True`.

Parameters:

| Name                     | Type                                   | Description  | Default            |
|--------------------------|--|--|--------------------|
| <code>delta</code>       | <code>bool</code>                      | if <code>True</code> , yield each chunk of text as it is received, if <code>False</code> (default), yield the full text up to the current point.   | <code>False</code> |
| <code>debounce_by</code> | <code>float</code>   <code>None</code> | by how much (if at all) to debounce/group the response chunks by. <code>None</code> means no debouncing. Debouncing is particularly important for long structured responses to reduce the overhead of performing validation as each token is received. | <code>0.1</code>   |

Source code in `pydantic_ai_slim/pydantic_ai/result.py`

```
158 async def stream_text(self, *, delta: bool = False, debounce_by: float | None = 0.1) -> AsyncIterator[str]:
159     """Stream the text result as an async iterable.
160
161     !!! note
162         This method will fail if the response is structured,
163         e.g. if [ 'is_structured' ][pydantic_ai.result.StreamedRunResult.is_structured] returns True.
164
165     !!! note
166         Result validators will NOT be called on the text result if delta=True.
167
168     Args:
169         delta: if True, yield each chunk of text as it is received, if False (default), yield the full text
170             up to the current point.
171         debounce_by: by how much (if at all) to debounce/group the response chunks by. None means no debouncing.
172             Debouncing is particularly important for long structured responses to reduce the overhead of
173             performing validation as each token is received.
174     """
175     with _logfire.span('response stream text') as lf_span:
176         if isinstance(self._stream_response, models.StreamStructuredResponse):
177             raise exceptions.UserError('stream_text() can only be used with text responses')
178         if delta:
179             async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
180                 async for _ in group_iter:
181                     yield ''.join(self._stream_response.get())
182                 final_delta = ''.join(self._stream_response.get(final=True))
183                 if final_delta:
184                     yield final_delta
185         else:
186             # a quick benchmark shows it's faster to build up a string with concat when we're
187             # yielding at each step
188             chunks: list[str] = []
189             combined = ''
190             async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
191                 async for _ in group_iter:
192                     new = False
193                     for chunk in self._stream_response.get():
194                         chunks.append(chunk)
195                         new = True
196                     if new:
197                         combined = await self._validate_text_result(''.join(chunks))
198                         yield combined
199             new = False
200             for chunk in self._stream_response.get(final=True):
201                 chunks.append(chunk)
202                 new = True
203             if new:
204                 combined = await self._validate_text_result(''.join(chunks))
205                 yield combined
206             lf_span.set_attribute('combined_text', combined)
207             self._marked_completed(text=combined)
208
```

stream\_structured async

```
stream_structured(
    *, debounce_by: float | None = 0.1
) -> AsyncIterator[tuple[ModelStructuredResponse, bool]]
```

Stream the response as an async iterable of Structured LLM Messages.

Note

This method will fail if the response is text, e.g. if `is_structured` returns `False`.

Parameters:

| Name                     | Type                                   | Description  | Default          |
|--------------------------|--|--|------------------|
| <code>debounce_by</code> | <code>float</code>   <code>None</code> | by how much (if at all) to debounce/group the response chunks by. <code>None</code> means no debouncing. Debouncing is particularly important for long structured responses to reduce the overhead of performing validation as each token is received. | <code>0.1</code> |

Returns:

| Type   | Description  |
|--|--|
| <code>AsyncIterator[tuple[ModelStructuredResponse, bool]]</code> | An async iterable of the structured response message and whether that is the last message. |

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
210 async def stream_structured(
211     self, *, debounce_by: float | None = 0.1
212 ) -> AsyncIterator[tuple[messages.ModelStructuredResponse, bool]]:
213     """Stream the response as an async iterable of Structured LLM Messages.
214
215     !!! note
216         This method will fail if the response is text,
217         e.g. if ['is_structured'][pydantic_ai.result.StreamedRunResult.is_structured] returns 'False'.
218
219     Args:
220         debounce_by: by how much (if at all) to debounce/group the response chunks by. 'None' means no debouncing.
221         Debouncing is particularly important for long structured responses to reduce the overhead of
222         performing validation as each token is received.
223
224     Returns:
225         An async iterable of the structured response message and whether that is the last message.
226     """
227     with _logfire.span('response stream structured') as lf_span:
228         if isinstance(self._stream_response, models.StreamTextResponse):
229             raise exceptions.UserError('stream_structured() can only be used with structured responses')
230         else:
231             # we should already have a message at this point, yield that first if it has any content
232             msg = self._stream_response.get()
233             if any(call.has_content() for call in msg.calls):
234                 yield msg, False
235             async with _utils.group_by_temporal(self._stream_response, debounce_by) as group_iter:
236                 async for _ in group_iter:
237                     msg = self._stream_response.get()
238                     if any(call.has_content() for call in msg.calls):
239                         yield msg, False
240             msg = self._stream_response.get(final=True)
241             yield msg, True
242             lf_span.set_attribute('structured_response', msg)
243             self._marked_completed(structured_message=msg)
```

**get\_data** async

```
get_data() -> ResultData
```

Stream the whole response, validate and return it.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
245 async def get_data(self) -> ResultData:
246     """Stream the whole response, validate and return it."""
247     async for _ in self._stream_response:
248         pass
249     if isinstance(self._stream_response, models.StreamTextResponse):
250         text = ''.join(self._stream_response.get(final=True))
251         text = await self._validate_text_result(text)
252         self._marked_completed(text=text)
253         return cast(ResultData, text)
254     else:
255         structured_message = self._stream_response.get(final=True)
256         self._marked_completed(structured_message=structured_message)
257         return await self.validate_structured_result(structured_message)
```

**is\_structured** property

```
is_structured: bool
```

Return whether the stream response contains structured data (as opposed to text).

**cost**

```
cost() -> Cost
```

Return the cost of the whole run.

#### Note

This won't return the full cost until the stream is finished.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
264 def cost(self) -> Cost:
265     """Return the cost of the whole run.
266
267     !!! note
268         This won't return the full cost until the stream is finished.
269     """
270     return self.cost_so_far + self._stream_response.cost()
```

**timestamp**

```
timestamp() -> datetime
```

Get the timestamp of the response.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
272 def timestamp(self) -> datetime:
273     """Get the timestamp of the response."""
274     return self._stream_response.timestamp()
```

**validate\_structured\_result** async

```
validate_structured_result(
    message: ModelStructuredResponse,
    *,
```

```
allow_partial: bool = False
) -> ResultData
```

Validate a structured result message.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
276 async def validate_structured_result(
277     self, message: messages.ModelStructuredResponse, *, allow_partial: bool = False
278 ) -> ResultData:
279     """Validate a structured result message."""
280     assert self._result_schema is not None, 'Expected _result_schema to not be None'
281     match = self._result_schema.find_tool(message)
282     if match is None:
283         raise exceptions.UnexpectedModelBehavior(
284             f'Invalid message, unable to find tool: {self._result_schema.tool_names()}'
285         )
286
287     call, result_tool = match
288     result_data = result_tool.validate(call, allow_partial=allow_partial, wrap_validation_errors=False)
289
290     for validator in self._result_validators:
291         result_data = await validator.validate(result_data, self._deps, 0, call)
292     return result_data
```

Cost dataclass

Cost of a request or run.

Responsibility for calculating costs is on the model used, PydanticAI simply sums the cost of requests.

You'll need to look up the documentation of the model you're using to convert "token count" costs to monetary costs.

99 Source code in pydantic\_ai\_slim/pydantic\_ai/result.py

```
28 @dataclass
29 class Cost:
30     """Cost of a request or run.
31
32     Responsibility for calculating costs is on the model used, PydanticAI simply sums the cost of requests.
33
34     You'll need to look up the documentation of the model you're using to convert "token count" costs to monetary costs.
35     """
36
37     request_tokens: int | None = None
38     """Tokens used in processing the request."""
39     response_tokens: int | None = None
40     """Tokens used in generating the response."""
41     total_tokens: int | None = None
42     """Total tokens used in the whole run, should generally be equal to 'request_tokens + response_tokens'."""
43     details: dict[str, int] | None = None
44     """Any extra details returned by the model."""
45
46     def __add__(self, other: Cost) -> Cost:
47         """Add two costs together.
48
49         This is provided so it's trivial to sum costs from multiple requests and runs.
50         """
51         counts: dict[str, int] = {}
52         for f in 'request_tokens', 'response_tokens', 'total_tokens':
53             self_value = getattr(self, f)
54             other_value = getattr(other, f)
55             if self_value is not None or other_value is not None:
56                 counts[f] = (self_value or 0) + (other_value or 0)
57
58         details = self.details.copy() if self.details is not None else None
59         if other.details is not None:
60             details = details or {}
61             for key, value in other.details.items():
62                 details[key] = details.get(key, 0) + value
63
64         return Cost(**counts, details=details or None)
```

**request\_tokens** class-attribute instance-attribute

```
request_tokens: int | None = None
```

Tokens used in processing the request.

**response\_tokens** class-attribute instance-attribute

```
response_tokens: int | None = None
```

Tokens used in generating the response.

**total\_tokens** class-attribute instance-attribute

```
total_tokens: int | None = None
```

Total tokens used in the whole run, should generally be equal to `request_tokens + response_tokens`.

**details** class-attribute instance-attribute

```
details: dict[str, int] | None = None
```

Any extra details returned by the model.

**\_\_add\_\_**

```
__add__(other: Cost) -> Cost
```

Add two costs together.

This is provided so it's trivial to sum costs from multiple requests and runs.



```
46 def __add__(self, other: Cost) -> Cost:
47     """Add two costs together.
48
49     This is provided so it's trivial to sum costs from multiple requests and runs.
50     """
51     counts: dict[str, int] = {}
52     for f in 'request_tokens', 'response_tokens', 'total_tokens':
53         self_value = getattr(self, f)
54         other_value = getattr(other, f)
55         if self_value is not None or other_value is not None:
56             counts[f] = (self_value or 0) + (other_value or 0)
57
58     details = self.details.copy() if self.details is not None else None
59     if other.details is not None:
60         details = details or {}
61         for key, value in other.details.items():
62             details[key] = details.get(key, 0) + value
63
64     return Cost(**counts, details=details or None)
```